

CSC B09H1 S 2017 Final Exam

Duration — 150 minutes

Aids allowed: Handwritten two-page
double-sided crib sheet

Student Number: _____

Utorid: _____

Last Name: _____ First Name: _____

Instructor: Bianca Schroeder

*Do **not** turn this page until you have received the signal to start.*

(Please fill out the identification section above and read the instructions below.)

Good Luck!

This final exam consists of 7 questions on 28 pages (including this one).
When you receive the signal to start, please make sure that your copy is complete.

Comments are not required except where indicated, although they may help us mark your answers. Note, that on the last pages of the exam we are providing some information from various Unix man pages that might be useful.

If you use any space for rough work, indicate clearly what you want marked.

1: _____/ 5

2: _____/10

3: _____/10

4: _____/10

5: _____/11

6: _____/ 9

7: _____/12

TOTAL: _____/67

Question 1. [5 MARKS]**Part (a)**

A parent process and a child share the same address space.

☐ True ☐ False

Part (b)

A process can call wait to wait for a process that it is not related to.

☐ True ☐ False

Part (c)

Signals can be used for communication between processes on different machines.

☐ True ☐ False

Part (d)

Internet routers guarantee reliable delivery of network packets.

☐ True ☐ False

Part (e)

TCP is more efficient, but less reliable than UDP.

☐ True ☐ False

Part (f)

A string's length (as indicated by strlen) and its size (as indicated by sizeof) are always equal.

☐ True ☐ False

Part (g)

Pipes can be used for communication between processes on different machines.

☐ True ☐ False

Part (h)

Sockets can be used for communication between processes running on the same machine.

☐ True ☐ False

Question 2. [10 MARKS]**Part (a)** [2 MARKS]

Considering the following piece of code, fill the table below with the values of the array elements after the code is done executing. Be careful with the difference between pointers and values, and pointer arithmetic.

```
int a[4] = {0, 1, 2, 3};
int b = 1;
int *p = a;
```

```
p = p + b;
b++;
*p += b;
p = p + b;
*p += 2;
p--;
*p *= 4;
p = p - b;
*p = p - a;
```

a[0]	a[1]	a[2]	a[3]

Part (b) [2 MARKS]

What is the output of the following program?

```
#include <stdio.h>

int func(int a, int *b, int *c) {
    a += 5;
    *b += a;
    c = b;
    return a;
}

int main() {
    int x = 5, y = 8, z = 3, t = 0;
    t = func(x, &y, &z);

    printf("x:%d    y:%d    z:%d    t:%d\n", x, y, z, t);
    return 0;
}
```

Answer:

Part (c) [4 MARKS]

There are several errors in the following program. Use the appropriate letter to label the lines of code where the corresponding error occurs.

P - dereferencing a pointer without having allocated memory for it.

D - deallocating memory that has already been deallocated.

H - deallocating memory that is not located on the heap.

M - memory leak.

```
int a = 0, b = 1, c = 2;
int *p, *q, *r;

p = malloc(sizeof(int));
*p = b;
q = &c;
p = &a;
*q = b;
q = malloc(sizeof(10));
*r = a + b + c;
free(p);
r = q;
p = malloc(sizeof(int));
*p = *r;
free(q);
free(r);
free(p);
```

Part (d) [2 MARKS]

What does the following piece of code print?

```
struct Point {
    int x;
    int y;
};

struct Point p1, p2;
struct Point *q1, *q2;

p1.x = 5;  p1.y = 10;
p2.x = 1;  p2.y = 10;

q1 = &p2;
q2 = &p1;
q1->x += 2;    q1->y += 7;
q2->x += 4;    q2->y += 3;

printf("P1(X,Y) = (%d, %d)\n", p1.x, p1.y);
printf("P2(X,Y) = (%d, %d)\n", p2.x, p2.y);
```

Answer:

Question 3. [10 MARKS]

Each example below contains an independent code fragment. In each case, there are variables x and y that are missing declaration statements. In the boxes to the right of the code, write those declaration statements so that the code fragment would compile and run without warnings or errors. If there is no declaration that could lead to a compilation without warnings or errors, write "ERROR". The first is done for you as an example.

Code Fragment	Declaration for x	Declaration for y
<pre>x = 10; y = 'A';</pre>	<pre>int x;</pre>	<pre>char y;</pre>
<pre>double length = 25; x = &length; y = &x;</pre>		
<pre>char *id[6]; x = id[3]; // some hidden code id[3] = "c3new"; y = *x[3];</pre>		
<pre>char *name = "John Tory"; x = &name; y = *(name + 3);</pre>		
<pre>struct node { int value; struct node *next; }; typedef struct node List; List *head; // some hidden code x = head->next; y.value = 14; y.next = x;</pre>		
<pre>char fun(char *str, int n) { return str[n]; } y = fun("hello", 1); x = &y;</pre>		

Question 4. [10 MARKS]

Some of the code fragments below have a problem. For each fragment indicate whether the code works as intended or there is an error (logical error, compile-time error/warning, or runtime error). Assume all programs are compiled using the *C99* standard. For this question, we will assume programs which do not terminate are errors as well. If there is an error in a fragment, explain **briefly** what is wrong in the box. For parts where there is an error you will only receive marks if you can correctly explain what is wrong. We have intentionally omitted the error checking of the system calls to simplify the examples. Do **not** report this as an error.

Some of the parts will use the following struct definition:

```
struct student {  
    int age;  
    char *name;  
};
```

```
char *s = "Hello";  
strcat(s, ", World!");
```

☐ Works as intended ☐ Error

```
int main(int argc, char **argv) {  
    char ch;  
    char *p = &ch;  
    ch = argv[argc-1][0];  
    printf("%c\n", p[0]);  
    return 0;  
}
```

☐ Works as intended ☐ Error

```
struct student hannah;  
strcpy(hannah.name, "Hannah");
```

☐ Works as intended ☐ Error

```
struct student hannah = NULL;  
// ... missing code ...  
if (hannah != NULL) {  
    hannah.age = 10;  
}
```

☐ Works as intended ☐ Error

```
// Increase the age of a student by amt.  
void increase_age(struct student s, int amt) {  
    s.age += amt;  
}
```

```
int main() {  
    struct student rob;  
    rob.age = 10;  
    increase_age(rob, 5);  
    printf("%d should be 15\n", rob.age);  
    return 0;  
}
```

☐ Works as intended ☐ Error


```
// Compute the sum of an array of integers
int compute_sum(int numbers[]) {
    int sum = 0;
    for (int i = 0; i < sizeof(numbers); i++) {
        sum += numbers[i];
    }
    return sum;
}
```

☐ Works as intended ☐ Error

```
int fd[2];
int result = fork();
pipe(fd);
if (result == 0) {
    close(fd[0]);
    write(fd[1], "cscb09", 7);
}
else {
    close(fd[1]);
    char buf[7];
    read(fd[0], buf, 7);
    printf("%s\n", buf);
}
exit(0);
```

☐ Works as intended ☐ Error

```
// Read all bytes from the file descriptors in 'fds' as characters,
// and print them. 'num_fds' is the number of file descriptors, and
// 'max_fd' is the value of the largest one.
void read_ints(int *fds, int num_fds, int max_fd) {
    char data;
    fd_set set;
    FD_ZERO(&set);
    for (int i = 0; i < num_fds; i++) {
        FD_SET(fds[i], &set);
    }

    while (select(max_fd + 1, &set, NULL, NULL, NULL) > 0) {
        for (int i = 0; i < num_fds; i++) {
            if (FD_ISSET(fds[i], &set)) {
                if (read(fds[i], &data, 1) > 0) {
                    printf("%c\n", data);
                }
            }
        }
    }
}
```

☐ Works as intended ☐ Error

```
struct node {
    int item;
    struct node *next;
};

// Compute the sum of the items in a linked list with the given head,
// but do not modify the list.
int sum(struct node *head) {
    int s = 0;
    while (head != NULL) {
        s += head->item;
        *head = *(head->next);
    }
    return s;
}
```

☐ Works as intended ☐ Error

```
// Remove the dots from word
char *word = "Ex.ampl.e";
char *result = malloc(strlen(word) + 1); // upper-limit if word has no dots
for (int i = 0; i < strlen(word); i++) {
    if (word[i] != '.') {
        strncat(result, word[i], 1);
    }
}
```

☐ Works as intended ☐ Error

Question 5. [11 MARKS]

Assume a linked list structure, which contains some information about a student, as follows:

```
typedef struct student {  
    int student_number;  
    char *last_name;  
    struct student *next;  
} Student;
```

Write a function that traverses a given student list and inserts a new student in alphabetical order (function prototype given below). If the new student has the same last name as another student from the list, insert in ascending order of the student number.

Note: The new student name passed to the function may be deallocated once the function exits, so make sure to create a new copy of the student name when inserting.

```
Student *insert_new_student(Student *list, int newStudentNo, char *newName) {
```

```
}
```

Question 6. [9 MARKS]

Consider the following program that runs to completion without errors:

```
int main() {
    int var = 1;
    int status;

    int r = fork();
    if(r == 0) {
        var++;
        r = fork();
        if(r == 0) { // process X
            var++;
            exit(var);
        } else { // process Y
            if(wait(&status) != -1) {
                if(WIFEXITED(status)) {
                    printf("A %d ", WEXITSTATUS(status));
                }
            }
            var += 2;
        }
    } else { // process Z
        printf("W %d ", var);
        if(wait(&status) != -1) {
            if(WIFEXITED(status)) {
                printf("B %d ", WEXITSTATUS(status));
            }
        }
    }
    printf("C %d ", var);
    return 0;
}
```

Part (a) [5 MARKS]

Write all the possible output orders for this program.

Part (b) [2 MARKS]

If all processes run to completion, can process X or process Y become an orphan? If yes, explain the sequence of events that would cause the process to become an orphan. If no, then explain what modification would need be made to the code so that a process might become an orphan. Clearly identify which process is the orphan.

Part (c) [2 MARKS]

If all processes run to completion, can process X or process Y become a zombie? If yes, explain the sequence of events that would cause the process to become a zombie. If no, then explain what modification would need be made to the code so that a process might become a zombie. Clearly identify which process is the zombie.

Question 7. [12 MARKS]

Implement a “parrot” server, which receives messages from clients and repeats each message back to the client who originally sent it. The server communicates with clients using *reliable* connections and uses the *I/O multiplexing* model.

Once a client connection is established, the server adds the new communication channel (socket) to a local client array to keep track of it (the code for initializing this is provided). Any new messages received from such a client will be repeated back to the client over the corresponding communication channel. The server does *not* timeout while waiting for client connections.

Notes:

- Use the API sheet. Syntax does not have to be perfect, but **do not write pseudocode**.
- Explain in comments all design decisions (size of queue for pending connections, etc.).
- Some code is already provided to guide you: extracting the port, initializing the array that will keep track of the sockets for clients, initializing file descriptor sets, etc. Have a look over the next page before starting.

```
int main(int argc, char *argv[]) {

    // Extract port form command line.
    int portno = strtol(argv[1], NULL, 10);

    // Initialize array of client socket descriptors.
    int i, clients[FD_SETSIZE];
    for(i = 0; i < FD_SETSIZE; i++)
        clients[i] = -1;

    // Fill server information.
    struct sockaddr_in serveraddr;
    memset(&serveraddr, 0, sizeof(serveraddr));
    serveraddr.sin_family = AF_INET;
    serveraddr.sin_addr.s_addr = htonl(INADDR_ANY);
    serveraddr.sin_port = htons((unsigned short) portno);

    // Add server operations for incoming client connections.
    int serverfd;
    ...
}
```

```
// Add serverfd to read set, set maxfd.  
fd_set rset, allset;  
FD_ZERO(&allset);  
FD_SET(serverfd, &allset);  
int maxfd = serverfd;
```

```
// Write the main server loop.  
for( ; ; ) {
```

```
} // end main server loop.
```

```
close(serverfd);  
return 0;  
} // end main.
```


Appendix

Useful structs

```
struct sockaddr_in {
    sa_family_t sin_family;
    unsigned short int sin_port;
    struct in_addr sin_addr;
    unsigned char pad[8]; /*Unused*/
}
```

Man page for `select`, `FD_CLR`, `FD_ISSET`, `FD_SET`, `FD_ZERO` - synchronous I/O multiplexing

SYNOPSIS

```
#include <sys/select.h>

int select(int nfds, fd_set *readfds, fd_set *writefds,
           fd_set *exceptfds, struct timeval *timeout);

void FD_CLR(int fd, fd_set *set);
int  FD_ISSET(int fd, fd_set *set);
void FD_SET(int fd, fd_set *set);
void FD_ZERO(fd_set *set);
```

`select()` allows a program to monitor multiple file descriptors, waiting until one or more of the file descriptors become "ready" for some class of I/O operation (e.g., input possible). A file descriptor is considered ready if it is possible to perform the corresponding I/O operation (e.g., `read(2)`) without blocking.

Three independent sets of file descriptors are watched. Those listed in `readfds` will be watched to see if characters become available for reading (more precisely, to see if a read will not block; in particular, a file descriptor is also ready on end-of-file), those in `writefds` will be watched to see if a write will not block, and those in `exceptfds` will be watched for exceptions. On exit, the sets are modified in place to indicate which file descriptors actually changed status. Each of the three file descriptor sets may be specified as `NULL` if no file descriptors are to be watched for the corresponding class of events.

Four macros are provided to manipulate the sets. `FD_ZERO()` clears a set. `FD_SET()` and `FD_CLR()` respectively add and remove a given file descriptor from a set. `FD_ISSET()` tests to see if a file descriptor is part of the set; this is useful after `select()` returns.

`nfds` is the highest-numbered file descriptor in any of the three sets, plus 1.

The `timeout` argument specifies the minimum interval that `select()` should block waiting for a file descriptor to become ready. (This interval will be rounded up to the system clock granularity, and kernel scheduling delays mean that the blocking interval may overrun by a small amount.) If both fields of the `timeval` structure are zero, then `select()` returns immediately. (This is useful for polling.) If `timeout` is `NULL` (no timeout), `select()` can block indefinitely.

RETURN VALUE

On success, `select()` returns the number of file descriptors contained in the three returned descriptor sets (that is, the total number of bits that are set in `readfds`, `writfds`, `exceptfds`) which may be zero if the timeout expires before anything interesting happens. On error, -1 is returned, and `errno` is set appropriately; the sets and timeout become undefined, so do not rely on their contents after an error.

Man page for socket

SYNOPSIS

```
#include <sys/types.h> /* See NOTES */
#include <sys/socket.h>
int socket(int domain, int type, int protocol);
```

DESCRIPTION

`socket()` creates an endpoint for communication and returns a descriptor. The domain argument specifies a communication domain; this selects the protocol family which will be used for communication. These families are defined in `sys/socket.h`. The currently understood formats include:

Name	Purpose	Man page
AF_UNIX	Local communication	unix(7)
AF_INET	IPv4 Internet protocols	ip(7)
AF_INET6	IPv6 Internet protocols	ipv6(7)
AF_IPX	IPX - Novell protocols	

The socket has the indicated type, which specifies the communication semantics. Currently defined types are:

SOCK_STREAM	Provides sequenced, reliable, two-way, connection-based byte streams. An out-of-band data transmission mechanism may be supported.
SOCK_DGRAM	Supports datagrams (connectionless, unreliable messages of a fixed maximum length).
SOCK_RAW	Provides raw network protocol access.

The protocol specifies a particular protocol to be used with the socket. Normally only a single protocol exists to support a particular socket type within a given protocol family, in which case protocol can be specified as 0. However, it is possible that many protocols may exist, in which case a particular protocol must be specified in this manner. The protocol number to use is specific to the 'communication domain' in which communication is to take place; see `protocols(5)`. See `getprotoent(3)` on how to map protocol name strings to protocol numbers.

Sockets of type `SOCK_STREAM` are full-duplex byte streams, similar to pipes. They do not preserve record boundaries. A stream socket must be in a connected state before any data may be sent or received on it. A connection to another socket is created with a `connect(2)` call. Once connected, data may be transferred using `read(2)` and `write(2)` calls or some variant of the `send(2)` and `recv(2)` calls. When a session has been completed a `close(2)` may be performed. Out-of-band data may also be transmitted as described in `send(2)` and received as described in `recv(2)`.

The communications protocols which implement a `SOCK_STREAM` ensure that data is not lost or duplicated. If a piece of data for which the peer protocol has buffer space cannot be successfully transmitted within a reasonable length of time, then the connection is considered to be dead. When `SO_KEEPALIVE` is enabled on the socket the protocol checks in a protocol-specific manner if the other end is still alive. A `SIGPIPE` signal is raised if a process sends or receives on a broken stream; this causes naive processes, which do not handle the signal, to exit. `SOCK_SEQPACKET` sockets employ the same system calls as `SOCK_STREAM` sockets. The only difference is that `read(2)` calls will return only the amount of data requested, and any data remaining in the arriving packet will be discarded. Also all message boundaries in incoming datagrams are preserved.

RETURN VALUE

On success, a file descriptor for the new socket is returned. On error, -1 is returned, and `errno` is set appropriately.

Man page for `connect` and `inet_pton`

SYNOPSIS

```
int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

DESCRIPTION

The `connect()` system call connects the socket referred to by the file descriptor `sockfd` to the address specified by `addr`. The `addrlen` argument specifies the size of `addr`. The format of the address in `addr` is determined by the address space of the socket `sockfd`; see `socket(2)` for further details. If the socket `sockfd` is of type `SOCK_DGRAM` then `addr` is the address to which datagrams are sent by default, and the only address from which datagrams are received. If the socket is of type `SOCK_STREAM` or `SOCK_SEQPACKET`, this call attempts to make a connection to the socket that is bound to the address specified by `addr`.

Generally, connection-based protocol sockets may successfully `connect()` only once; connectionless protocol sockets may use `connect()` multiple times to change their association. Connectionless sockets may dissolve the association by connecting to an address with the `sa_family` member of `sockaddr` set to `AF_UNSPEC` (supported on Linux since kernel 2.2).

RETURN VALUE

If the connection or binding succeeds, zero is returned. On error, -1 is returned, and `errno` is set appropriately.

SYNOPSIS

```
#include <arpa/inet.h>
```

```
int inet_pton(int af, const char *src, void *dst);
```

DESCRIPTION

This function converts the character string `src` into a network address structure in the `af` address family, then copies the network address structure to `dst`. The `af` argument must be either `AF_INET` or `AF_INET6`.

`inet_pton()` returns 1 on success (network address was successfully converted). 0 is returned if `src` does not contain a character string representing a valid network address in the specified address family. If `af` does not contain a valid address family, -1 is returned and `errno` is set to `EAFNOSUPPORT`.

Man page for listen**SYNOPSIS**

```
int listen(int sockfd, int backlog);
```

DESCRIPTION

listen() marks the socket referred to by sockfd as a passive socket, that is, as a socket that will be used to accept incoming connection requests using accept(2).

The sockfd argument is a file descriptor that refers to a socket of type SOCK_STREAM or SOCK_SEQPACKET.

The backlog argument defines the maximum length to which the queue of pending connections for sockfd may grow. If a connection request arrives when the queue is full, the client may receive an error with an indication of ECONNREFUSED or, if the underlying protocol supports retransmission, the request may be ignored so that a later reattempt at connection succeeds.

RETURN VALUE

On success, zero is returned. On error, -1 is returned, and errno is set appropriately.

Man page for bind

SYNOPSIS

```
int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

DESCRIPTION

When a socket is created with `socket(2)`, it exists in a name space (address family) but has no address assigned to it. `bind()` assigns the address specified by `addr` to the socket referred to by the file descriptor `sockfd`. `addrlen` specifies the size, in bytes, of the address structure pointed to by `addr`. Traditionally, this operation is called assigning a name to a socket.

It is normally necessary to assign a local address using `bind()` before a `SOCK_STREAM` socket may receive connections (see `accept(2)`).

The rules used in name binding vary between address families.

The actual structure passed for the `addr` argument will depend on the address family. The `sockaddr` structure is defined as something like:

```
struct sockaddr {
    sa_family_t sa_family;
    char        sa_data[14];
}
```

The only purpose of this structure is to cast the structure pointer passed in `addr` in order to avoid compiler warnings. See `EXAMPLE` below.

RETURN VALUE

On success, zero is returned. On error, `-1` is returned, and `errno` is set appropriately.

Man page for htonl, htons, ntohl, ntohs**SYNOPSIS**

```
#include <arpa/inet.h>
uint32_t htonl(uint32_t hostlong);
uint16_t htons(uint16_t hostshort);
uint32_t ntohl(uint32_t netlong);
uint16_t ntohs(uint16_t netshort);
```

DESCRIPTION

The `htonl()` function converts the unsigned integer `hostlong` from host byte order to network byte order. The `htons()` function converts the unsigned short integer `hostshort` from host byte order to network byte order.

The `ntohl()` function converts the unsigned integer `netlong` from network byte order to host byte order.

The `ntohs()` function converts the unsigned short integer `netshort` from network byte order to host byte order.

On the i386 the host byte order is Least Significant Byte first, whereas the network byte order, as used on the Internet, is Most Significant Byte first.

Man page for fread**SYNOPSIS**

```
#include <stdio.h>
size_t fread(void *ptr, size_t size, size_t nmemb, FILE *stream );

size_t fwrite(const void *ptr, size_t size, size_t nmemb, FILE *stream);
```

DESCRIPTION

The function `fread()` reads `nmemb` elements of data, each `size` bytes long, from the stream pointed to by `stream`, storing them at the location given by `ptr`. The function `fwrite()` writes `nmemb` elements of data, each `size` bytes long, to the stream pointed to by `stream`, obtaining them from the location given by `ptr`.

RETURN VALUE

On success, `fread()` and `fwrite()` return the number of items read or written. This number equals the number of bytes transferred only when `size` is 1. If an error occurs, or the end of the file is reached, the return value is a short item count (or zero). `fread()` does not distinguish between end-of-file and error, and callers must use `feof(3)` and `ferror(3)` to determine which occurred.

Man page for read**SYNOPSIS**

```
#include <unistd.h>
ssize_t read(int fd, void *buf, size_t count);
```

DESCRIPTION

`read()` attempts to read up to `count` bytes from file descriptor `fd` into the buffer starting at `buf`. On files that support seeking, the read operation commences at the current file offset, and the file offset is incremented by the number of bytes read. If the current file offset is at or past the end of file, no bytes are read, and `read()` returns zero.

If `count` is zero, `read()` may detect the errors described below. In the absence of any errors, or if `read()` does not check for errors, a `read()` with a count of 0 returns zero and has no other effects.

If `count` is greater than `SSIZE_MAX`, the result is unspecified.

RETURN VALUE

On success, the number of bytes read is returned (zero indicates end of file), and the file position is advanced by this number. It is not an error if this number is smaller than the number of bytes requested; this may happen for example because fewer bytes are actually available right now (maybe because we were close to end-of-file, or because we are reading from a pipe, or from a terminal), or because `read()` was interrupted by a signal. On error, -1 is returned, and `errno` is set appropriately. In this case it is left unspecified whether the file position (if any) changes.

Man page for write

SYNOPSIS

```
#include <unistd.h>
ssize_t write(int fd, const void *buf, size_t count);
```

DESCRIPTION

`write()` writes up to `count` bytes from the buffer pointed `buf` to the file referred to by the file descriptor `fd`. The number of bytes written may be less than `count` if, for example, there is insufficient space on the underlying physical medium, or the `RLIMIT_FSIZE` resource limit is encountered (see `setrlimit(2)`), or the call was interrupted by a signal handler after having written less than `count` bytes. (See also `pipe(7)`.)

For a seekable file (i.e., one to which `lseek(2)` may be applied, for example, a regular file) writing takes place at the current file offset, and the file offset is incremented by the number of bytes actually written. If the file was `open(2)`ed with `O_APPEND`, the file offset is first set to the end of the file before writing. The adjustment of the file offset and the write operation are performed as an atomic step.

POSIX requires that a `read(2)` which can be proved to occur after a `write()` has returned returns the new data. Note that not all file systems are POSIX conforming.

RETURN VALUE

On success, the number of bytes written is returned (zero indicates nothing was written). On error, -1 is returned, and `errno` is set appropriately. If `count` is zero and `fd` refers to a regular file, then `write()` may return a failure status if one of the errors below is detected. If no errors are detected, 0 will be returned without causing any other effect. If `count` is zero and `fd` refers to a file other than a regular file, the results are not specified.

Man page for strcpy and strncpy

SYNOPSIS

```
#include <string.h>
```

```
char *strcpy(char *dest, const char *src);
```

```
char *strncpy(char *dest, const char *src, size_t n);
```

DESCRIPTION

The `strcpy()` function copies the string pointed to by `src`, including the terminating null byte, to the buffer pointed to by `dest`. The strings may not overlap, and the destination string `dest` must be large enough to receive the copy.

The `strncpy()` function is similar, except that at most `n` bytes of `src` are copied.

Warning: If there is no null byte among the first `n` bytes of `src`, the string placed in `dest` will not be null terminated.

If the length of `src` is less than `n`, `strncpy()` pads the remainder of `dest` with null bytes.

RETURN VALUE

The `strcpy()` and `strncpy()` functions return a pointer to the destination string `dest`.

Man page for strcat and strncat

SYNOPSIS

```
#include <string.h>
```

```
char *strcat(char *dest, const char *src);
```

```
char *strncat(char *dest, const char *src, size_t n);
```

DESCRIPTION

The `strcat()` function appends the `src` string to the `dest` string, overwriting the terminating null

The `strncat()` function is similar, except that

*

it will use at most `n` bytes from `src`; and

*

`src` does not need to be null-terminated if it contains `n` or more bytes.

As with `strcat()`, the resulting string in `dest` is always null-terminated.

If `src` contains `n` or more bytes, `strncat()` writes `n+1` bytes to `dest` (`n` from `src` plus the terminating null byte).

Man page for `strlen`

SYNOPSIS

```
#include <string.h>
size_t strlen(const char *s);
```

DESCRIPTION

The `strlen()` function calculates the length of the string `s`, excluding the terminating null byte.

RETURN VALUE

The `strlen()` function returns the number of bytes in the string `s`.

Man page for `fopen`

SYNOPSIS

```
#include <stdio.h>

FILE *fopen(const char *path, const char *mode);
```

DESCRIPTION

The `fopen()` function opens the file whose name is the string pointed to by `path` and associates a stream with it.

The argument `mode` points to a string beginning with one of the following sequences (possibly followed by additional characters, as described below):

`r`

Open text file for reading. The stream is positioned at the beginning of the file.

`r+`

Open for reading and writing. The stream is positioned at the beginning of the file.

`w`

Truncate file to zero length or create text file for writing. The stream is positioned

at the beginning of the file.

RETURN VALUE

Upon successful completion `fopen()` returns a FILE pointer.
Otherwise, NULL is returned and `errno` is set to indicate the error.

Man page for fork

SYNOPSIS

```
#include <unistd.h>
pid_t fork(void);
```

DESCRIPTION

`fork()` creates a new process by duplicating the calling process. The new process, referred to as the child, is an exact duplicate of the calling process, referred to as the parent, except for the following points:

- * The child has its own unique process ID, and this PID does not match the ID of any existing process group (`setpgid(2)`).
- * The child's parent process ID is the same as the parent's process ID.

RETURN VALUE

On success, the PID of the child process is returned in the parent, and 0 is returned in the child. On failure, -1 is returned in the parent, no child process is created, and `errno` is set appropriately.

Man page for scanf

SYNOPSIS

```
#include <stdio.h>

int scanf(const char *format, ...);
int fscanf(FILE *stream, const char *format, ...);
```

DESCRIPTION

The `scanf()` family of functions scans input according to format as described below. This format may contain conversion specifications; the results from such conversions, if any, are stored in the locations pointed to by the pointer arguments that follow format. Each pointer argument must be of a type that is appropriate for the value returned by the corresponding conversion specification.

If the number of conversion specifications in format exceeds the number of pointer arguments, the results are undefined. If the number of pointer arguments exceeds the number of conversion specifications, then the excess pointer arguments are evaluated, but are otherwise ignored.

The `scanf()` function reads input from the standard input stream `stdin`, `fscanf()` reads input from the stream pointer `stream`, and `sscanf()` reads its input from the character string pointed to by `str`.

The format string consists of a sequence of directives which describe how to process the sequence of input characters. If processing of a directive fails, no further input is read, and `scanf()` returns. A "failure" can be either of the following: input failure, meaning that input characters were unavailable, or matching failure, meaning that the input was inappropriate.

Man page for `malloc`

SYNOPSIS

```
#include <stdlib.h>
void *malloc(size_t size);
```

DESCRIPTION

The `malloc()` function allocates `size` bytes and returns a pointer to the allocated memory. The memory is not initialized. If `size` is 0, then `malloc()` returns either `NULL`, or a unique pointer value that can later be successfully passed to `free()`.

RETURN VALUE

The `malloc()` function return a pointer to the allocated memory that is suitably aligned for any kind of variable. On error, these functions return `NULL`. `NULL` may also be returned by a successful call to `malloc()` with a size of zero, or by a successful call to `calloc()` with `nmemb` or `size` equal to zero.

Man page for `wait`

Name

`wait` - wait for process to change state

Synopsis

```
#include <sys/types.h>
#include <sys/wait.h>
```

```
pid_t wait(int *status);
```

All of these system calls are used to wait for state changes in a child of the calling process, and obtain information about the child whose state has changed. A state change is considered to be: the child terminated; the child was stopped by a signal; or the child was resumed by a signal. In the case of a terminated child, performing a `wait` allows the system to release the resources associated with the child; if a `wait` is not performed, then

the terminated child remains in a "zombie" state.

If a child has already changed state, then these calls return immediately. Otherwise they block until either a child changes state or a signal handler interrupts the call (assuming that system calls are not automatically restarted using the SA_RESTART flag of `sigaction(2)`). In the remainder of this page, a child whose state has changed and which has not yet been waited upon by one of these system calls is termed waitable.

The `wait()` system call suspends execution of the calling process until one of its children terminates.

`WIFEXITED(status)`

returns true if the child terminated normally, that is, by calling `exit(3)` or `_exit(2)`, or by returning from `main()`.

`WEXITSTATUS(status)`

returns the exit status of the child. This consists of the least significant 8 bits of the status argument that the child specified in a call to `exit(3)` or `_exit(2)` or as the argument for a return statement in `main()`. This macro should only be employed if `WIFEXITED` returned true.

RETURN VALUE

`wait()`: on success, returns the process ID of the terminated child; on error, -1 is returned.

